```python
import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import models, transforms
from PIL import Image
import matplotlib.pyplot as plt
import numpy as np
import random

# 1. SETUP & DEVICE
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"□ Launching Project Vibe on: {device}")

# 2. VGG19 SAFETY PATCH (No Inplace)
vgg =
models.vgg19(weights=models.VGG19_Weights.DEFAULT).features.to(device)
.eval()
for m in vgg.modules():
    if isinstance(m, nn.ReLU): m.inplace = False
for p in vgg.parameters(): p.requires_grad = False

# 3. KAGGLE PATH FINDER
base_search = '/kaggle/input'
img_dir = None
for root, dirs, files in os.walk(base_search):
    if len([f for f in files if f.endswith('.jpg')]) > 50:
        img_dir = root
        break
if not img_dir: raise FileNotFoundError("□ Add CelebA dataset to
sidebar!")

# 4. LOAD & PREP IMAGES
all_imgs = [os.path.join(img_dir, f) for f in os.listdir(img_dir) if
f.endswith('.jpg')]
p1, p2 = random.sample(all_imgs, 2)
tf = transforms.Compose([transforms.Resize((256, 256)),
transforms.ToTensor(), transforms.Normalize([0.5]*3, [0.5]*3)])

sub1 = tf(Image.open(p1).convert('RGB')).unsqueeze(0).to(device)
sub2 = tf(Image.open(p2).convert('RGB')).unsqueeze(0).to(device)

# 5. MASK & UTILS
def make_face_mask(h=256, w=256):
    Y, X = torch.meshgrid(torch.linspace(0, 1, h), torch.linspace(0,
1, w), indexing='ij')
    dist = ((Y - 0.42)/0.38)**2 + ((X - 0.5)/0.28)**2
    return torch.clamp(1.0 - dist, 0.0, 1.0).to(device)
```

```
face_mask = make_face_mask()

class HierarchicalDSL(nn.Module):
    def __init__(self):
        super().__init__()
        sx = torch.tensor([[-1.,0.,1.],[-2.,0.,2.],[-
1.,0.,1.]]).view(1,1,3,3)
        sy = torch.tensor([[-1.,-2.,-1.],[0.,0.,0.],
[1.,2.,1.]]).view(1,1,3,3)
        self.register_buffer('sx', sx); self.register_buffer('sy', sy)
    def forward(self, pred, ref):
        p, r = pred.mean(1, keepdim=True), ref.mean(1, keepdim=True)
        px, py = F.conv2d(p, self.sx, padding=1), F.conv2d(p, self.sy,
padding=1)
        rx, ry = F.conv2d(r, self.sx, padding=1), F.conv2d(r, self.sy,
padding=1)
        pm, rm = torch.sqrt(px**2 + py**2 + 1e-6), torch.sqrt(rx**2 +
ry**2 + 1e-6)
        cos = (px/pm)*(rx/rm) + (py/pm)*(ry/rm)
        return ((1.0 - cos) * rm).mean()

def gram(f):
    b, c, h, w = f.shape
    F_ = f.view(b*c, h*w)
    return torch.mm(F_, F_.t()).div(b * c * h * w)

# 6. PRE-COMPUTE TARGET FEATURES
LAYER_MAP = {'3':'s1', '8':'s2', '17':'s3', '26':'s4', '21':'c1'}
with torch.no_grad():
    s2_feats = {}; x = sub2
    for name, layer in vgg._modules.items():
        x = layer(x)
        if name in LAYER_MAP: s2_feats[LAYER_MAP[name]] = x.detach()

# 7. INITIALIZE TARGET
target = (0.65 * sub1 + 0.35 * sub2).clone().requires_grad_(True)
hdsl = HierarchicalDSL().to(device)
optimizer = optim.LBFGS([target], lr=0.8, max_iter=20,
line_search_fn='strong_wolfe')

# 8. OPTIMIZATION
def closure():
    optimizer.zero_grad()
    target.data.clamp_(-1.0, 1.0)
    x_t = target; c_loss = s_loss = torch.tensor(0.0, device=device)
    for name, layer in vgg._modules.items():
        x_t = layer(x_t)
        if name in LAYER_MAP:
            key = LAYER_MAP[name]
            if key == 'c1': c_loss = F.mse_loss(x_t, s2_feats['c1'])
```

```
            else:
                w = 1.0 if key in ('s1', 's2') else 0.3
                s_loss += w * F.mse_loss(gram(x_t),
gram(s2_feats[key]))

    h_loss = hdsl(target, sub1)
    bg_mask = (1.0 - face_mask).unsqueeze(0).unsqueeze(0)
    bg_loss = (bg_mask * (target - sub1)**2).mean()

    # Weights rebalanced for presentation quality
    total = (3.0 * c_loss) + (2e4 * s_loss) + (1200.0 * h_loss) +
(500.0 * bg_loss) + (3.0 * (torch.abs(target[:,:,:,1:] -
target[:,:,:,:-1]).mean() + torch.abs(target[:,:,1:,:] - target[:,:,:-
1,:]).mean()))
    total.backward(); return total

print("□ Running Final Fusion... This will take ~2 mins.")
for i in range(15):
    optimizer.step(closure)
    print(f"Batch {i+1}/15 Complete")

# 9. OUTPUT
def denorm(t): return
(t.cpu().detach().squeeze().permute(1,2,0)*0.5+0.5).clamp(0,1)
plt.figure(figsize=(18, 5))
titles = ["Sub 1 (Structure)", "Sub 2 (Style)", "□ Final Result",
"Face Mask"]
imgs = [denorm(sub1), denorm(sub2), denorm(target),
face_mask.cpu().numpy()]
for i in range(4):
    plt.subplot(1, 4, i+1); plt.imshow(imgs[i], cmap='hot' if i==3
else None)
    plt.title(titles[i]); plt.axis('off')
plt.savefig('project_vibe_final.png', bbox_inches='tight'); plt.show()
```

```
□ Launching Project Vibe on: cuda
Downloading: "https://download.pytorch.org/models/vgg19-dcbb9e9d.pth"
to /root/.cache/torch/hub/checkpoints/vgg19-dcbb9e9d.pth

100%|██████████| 548M/548M [00:02<00:00, 209MB/s]

□ Running Final Fusion... This will take ~2 mins.
Batch 1/15 Complete
Batch 2/15 Complete
Batch 3/15 Complete
Batch 4/15 Complete
Batch 5/15 Complete
Batch 6/15 Complete
Batch 7/15 Complete
Batch 8/15 Complete
```

```
Batch 9/15 Complete
Batch 10/15 Complete
Batch 11/15 Complete
Batch 12/15 Complete
Batch 13/15 Complete
Batch 14/15 Complete
Batch 15/15 Complete

/tmp/ipykernel_24/1671283433.py:114: UserWarning: Glyph 127919 (\
N{DIRECT HIT}) missing from font(s) DejaVu Sans.
  plt.savefig('project_vibe_final.png', bbox_inches='tight');
plt.show()
/usr/local/lib/python3.12/dist-packages/IPython/core/pylabtools.py:151
: UserWarning: Glyph 127919 (\N{DIRECT HIT}) missing from font(s)
DejaVu Sans.
  fig.canvas.print_figure(bytes_io, **kw)
```

```
import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import models, transforms
from PIL import Image
import matplotlib.pyplot as plt
import numpy as np
import random

# 1. SETUP & SAFETY PATCH
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
vgg =
models.vgg19(weights=models.VGG19_Weights.DEFAULT).features.to(device)
.eval()
for m in vgg.modules():
    if isinstance(m, nn.ReLU): m.inplace = False
for p in vgg.parameters(): p.requires_grad = False

# 2. KAGGLE PATH FINDER
img_dir = None
```

```
for root, dirs, files in os.walk('/kaggle/input'):
    if len([f for f in files if f.endswith('.jpg')]) > 50:
        img_dir = root; break
if not img_dir: raise FileNotFoundError("□ Add CelebA to sidebar!")

# 3. LOAD IMAGES
all_imgs = [os.path.join(img_dir, f) for f in os.listdir(img_dir) if
f.endswith('.jpg')]
p1, p2 = random.sample(all_imgs, 2)
tf = transforms.Compose([transforms.Resize((256, 256)),
transforms.ToTensor(), transforms.Normalize([0.5]*3, [0.5]*3)])
sub1 = tf(Image.open(p1).convert('RGB')).unsqueeze(0).to(device)
sub2 = tf(Image.open(p2).convert('RGB')).unsqueeze(0).to(device)

# 4. ZONED FACE MASK (v4 logic)
def make_zoned_masks(h=256, w=256):
    Y, X = torch.meshgrid(torch.linspace(0, 1, h), torch.linspace(0,
1, w), indexing='ij')
    dist = ((Y - 0.42)/0.38)**2 + ((X - 0.5)/0.28)**2
    core = torch.clamp(1.0 - dist / 0.25, 0.0, 1.0)
    full = torch.clamp(1.0 - dist, 0.0, 1.0)
    bg = 1.0 - full
    return core.to(device), full.to(device), bg.to(device)

core_mask, full_mask, bg_mask = make_zoned_masks()

# 5. UTILS
class HierarchicalDSL(nn.Module):
    def __init__(self):
        super().__init__()
        sx = torch.tensor([[-1.,0.,1.],[-2.,0.,2.],[-
1.,0.,1.]]).view(1,1,3,3)
        sy = torch.tensor([[-1.,-2.,-1.],[0.,0.,0.],
[1.,2.,1.]]).view(1,1,3,3)
        self.register_buffer('sx', sx); self.register_buffer('sy', sy)
    def forward(self, pred, ref):
        p, r = pred.mean(1, keepdim=True), ref.mean(1, keepdim=True)
        px, py = F.conv2d(p, self.sx, padding=1), F.conv2d(p, self.sy,
padding=1)
        rx, ry = F.conv2d(r, self.sx, padding=1), F.conv2d(r, self.sy,
padding=1)
        pm, rm = torch.sqrt(px**2 + py**2 + 1e-6), torch.sqrt(rx**2 +
ry**2 + 1e-6)
        cos = (px/pm)*(rx/rm) + (py/pm)*(ry/rm)
        return ((1.0 - cos) * rm).mean()

def gram(f):
    b, c, h, w = f.shape
    return torch.mm(f.view(b*c, h*w), f.view(b*c, h*w).t()).div(b * c
* h * w)
```

```
# 6. PRE-COMPUTE
LAYER_MAP = {'3':'s1', '8':'s2', '14':'s3', '17':'s4', '21':'c1',
'26':'c2', '30':'c3'}
with torch.no_grad():
    s2_feats = {}; x = sub2
    for name, layer in vgg._modules.items():
        x = layer(x)
        if name in LAYER_MAP: s2_feats[LAYER_MAP[name]] = x.detach()

target = (0.65 * sub1 + 0.35 * sub2).clone().requires_grad_(True)
hdsl = HierarchicalDSL().to(device)
optimizer = optim.LBFGS([target], lr=0.8, max_iter=20,
line_search_fn='strong_wolfe')

# 7. OPTIMIZATION
def closure():
    optimizer.zero_grad()
    target.data.clamp_(-1.0, 1.0)
    x_t = target; c_loss = s_loss = torch.tensor(0.0, device=device)
    for name, layer in vgg._modules.items():
        x_t = layer(x_t)
        if name in LAYER_MAP:
            key = LAYER_MAP[name]
            if key in ['c1', 'c2', 'c3']:
                w = {'c1': 1.0, 'c2': 0.5, 'c3': 0.3}[key]
                c_loss += w * F.mse_loss(x_t, s2_feats[key])
            else:
                w = {'s1': 0.2, 's2': 0.2, 's3': 0.8, 's4': 0.8}[key]
                s_loss += w * F.mse_loss(gram(x_t),
gram(s2_feats[key]))

    total = (6.0 * c_loss) + (1e4 * s_loss) + (1000.0 * hdsl(target,
sub1)) + \
            (600.0 * (bg_mask.unsqueeze(0).unsqueeze(0) * (target -
sub1)**2).mean()) + \
            (2.0 * (torch.abs(target[:,:,:,1:] - target[:,:,:,:-
1]).mean() + torch.abs(target[:,:,1:,:] - target[:,:,:-1,:]).mean()))
    total.backward(); return total

print("□ CDSM v4 — Processing...")
for i in range(15):
    optimizer.step(closure)
    print(f"Batch {i+1}/15 done")

# 8. DISPLAY
def denorm(t): return
(t.cpu().detach().squeeze().permute(1,2,0)*0.5+0.5).clamp(0,1)
plt.figure(figsize=(20, 10))
imgs = [denorm(sub1), denorm(sub2), denorm(target)]
```

```
titles = ["Sub 1 (Structure)", "Sub 2 (Identity)", "□ v4 Result"]
for i in range(3):
    plt.subplot(2, 3, i+1); plt.imshow(imgs[i]); plt.title(titles[i]);
plt.axis('off')

# Zone Map Visualization
zone_img = np.zeros((256, 256, 3))
zone_img[:,:,0] = bg_mask.cpu().numpy() # Red = BG
zone_img[:,:,2] = core_mask.cpu().numpy() # Blue = Core
plt.subplot(2, 3, 4); plt.imshow(zone_img); plt.title("Spatial Zone
Map"); plt.axis('off')
plt.show()
```

```
□ CDSM v4 — Processing...
Batch 1/15 done
Batch 2/15 done
Batch 3/15 done
Batch 4/15 done
Batch 5/15 done
Batch 6/15 done
Batch 7/15 done
Batch 8/15 done
Batch 9/15 done
Batch 10/15 done
Batch 11/15 done
Batch 12/15 done
Batch 13/15 done
Batch 14/15 done
Batch 15/15 done
```

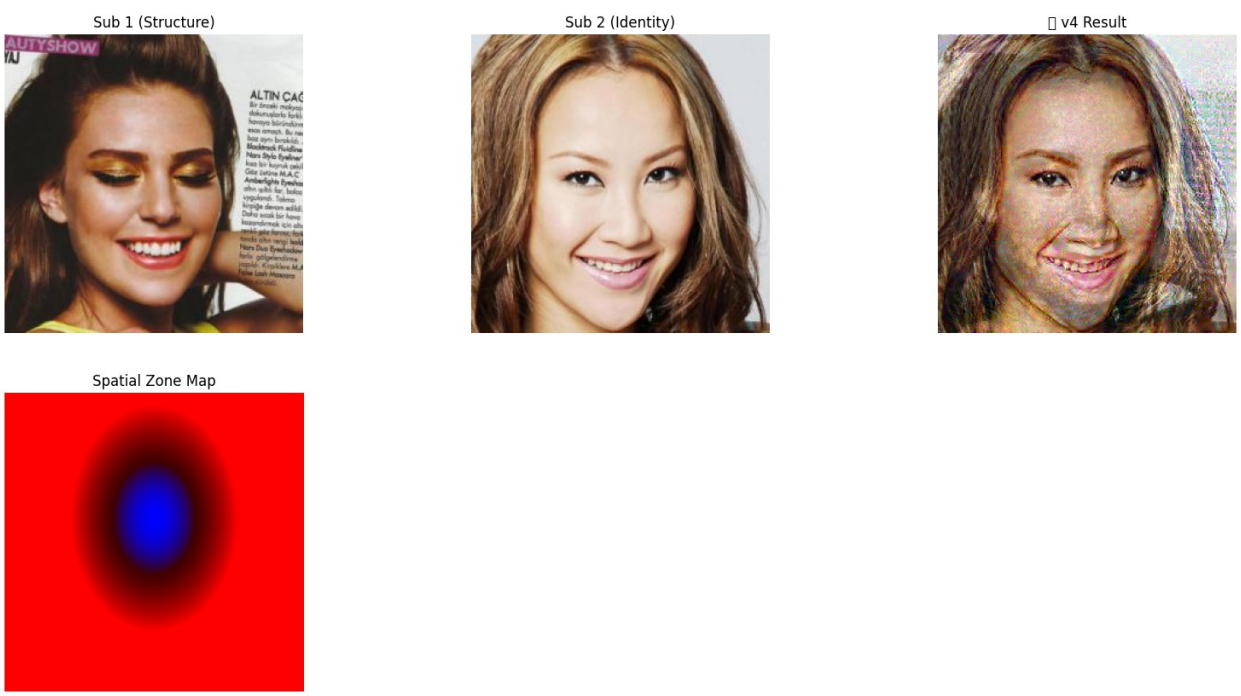

```python
import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import models, transforms
from PIL import Image
import matplotlib.pyplot as plt
import numpy as np
import random

# ═══════════════════════════════════════════════════════
# 1. ENVIRONMENT & MODEL SETUP
# ═══════════════════════════════════════════════════════
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"□ Project Vibe Final Model active on: {device}")

vgg =
models.vgg19(weights=models.VGG19_Weights.DEFAULT).features.to(device)
.eval()
for m in vgg.modules():
    if isinstance(m, nn.ReLU): m.inplace = False
for p in vgg.parameters(): p.requires_grad = False

# ═══════════════════════════════════════════════════════
# 2. DATASET PATH FINDER
# ═══════════════════════════════════════════════════════
img_dir = None
for root, dirs, files in os.walk('/kaggle/input'):
    if len([f for f in files if f.endswith('.jpg')]) > 50:
        img_dir = root; break
if not img_dir: raise FileNotFoundError("□ Please add CelebA dataset
to the sidebar!")

# ═══════════════════════════════════════════════════════
# 3. ZONED SPATIAL MASK & H-DSL
# ═══════════════════════════════════════════════════════
def make_zoned_masks(h=256, w=256):
    Y, X = torch.meshgrid(torch.linspace(0, 1, h), torch.linspace(0,
1, w), indexing='ij')
    dist = ((Y - 0.42)/0.38)**2 + ((X - 0.5)/0.28)**2
    core = torch.clamp(1.0 - dist / 0.20, 0.0, 1.0) # Tight identity
core
    full = torch.clamp(1.0 - dist, 0.0, 1.0)        # Texture mid-zone
    bg = 1.0 - full                                 # Background
anchor
    return core.to(device), full.to(device), bg.to(device)

core_mask, full_mask, bg_mask = make_zoned_masks()
```

```
class HierarchicalDSL(nn.Module):
    def __init__(self):
        super().__init__()
        sx = torch.tensor([[-1.,0.,1.],[-2.,0.,2.],[-
1.,0.,1.]]).view(1,1,3,3)
        sy = torch.tensor([[-1.,-2.,-1.],[0.,0.,0.],
[1.,2.,1.]]).view(1,1,3,3)
        self.register_buffer('sx', sx); self.register_buffer('sy', sy)
    def forward(self, pred, ref):
        p, r = pred.mean(1, keepdim=True), ref.mean(1, keepdim=True)
        px, py = F.conv2d(p, self.sx, padding=1), F.conv2d(p, self.sy,
padding=1)
        rx, ry = F.conv2d(r, self.sx, padding=1), F.conv2d(r, self.sy,
padding=1)
        pm, rm = torch.sqrt(px**2 + py**2 + 1e-6), torch.sqrt(rx**2 +
ry**2 + 1e-6)
        cos = (px/pm)*(rx/rm) + (py/pm)*(ry/rm)
        return ((1.0 - cos) * rm).mean()

def gram(f):
    b, c, h, w = f.shape
    return torch.mm(f.view(b*c, h*w), f.view(b*c, h*w).t()).div(b * c
* h * w)

# ═══════════════════════════════════════════════════
# 4. INITIALIZATION & PRE-COMPUTE
# ═══════════════════════════════════════════════════
all_imgs = [os.path.join(img_dir, f) for f in os.listdir(img_dir) if
f.endswith('.jpg')]
p1, p2 = random.sample(all_imgs, 2)
tf = transforms.Compose([transforms.Resize((256, 256)),
transforms.ToTensor(), transforms.Normalize([0.5]*3, [0.5]*3)])
sub1 = tf(Image.open(p1).convert('RGB')).unsqueeze(0).to(device)
sub2 = tf(Image.open(p2).convert('RGB')).unsqueeze(0).to(device)

LAYER_MAP = {'3':'s1', '8':'s2', '14':'s3', '17':'s4', '21':'c1',
'26':'c2', '30':'c3'}
with torch.no_grad():
    s2_feats = {}; x = sub2
    for name, layer in vgg._modules.items():
        x = layer(x)
        if name in LAYER_MAP: s2_feats[LAYER_MAP[name]] = x.detach()

target = (0.65 * sub1 + 0.35 * sub2).clone().requires_grad_(True)
hdsl = HierarchicalDSL().to(device)
optimizer = optim.LBFGS([target], lr=0.8, max_iter=25,
line_search_fn='strong_wolfe')

# ═══════════════════════════════════════════════════
# 5. OPTIMIZATION LOOP
```

```
# ═══════════════════════════════════════════════════
def closure():
    optimizer.zero_grad()
    target.data.clamp_(-1.0, 1.0)
    x_t = target; c_loss = s_loss = torch.tensor(0.0, device=device)

    for name, layer in vgg._modules.items():
        x_t = layer(x_t)
        if name in LAYER_MAP:
            key = LAYER_MAP[name]
            if key in ['c1', 'c2', 'c3']:
                w = {'c1': 1.0, 'c2': 0.8, 'c3': 0.5}[key]
                c_loss += w * F.mse_loss(x_t, s2_feats[key])
            else:
                w = {'s1': 0.1, 's2': 0.1, 's3': 0.9, 's4': 0.9}[key]
                s_loss += w * F.mse_loss(gram(x_t),
gram(s2_feats[key]))

    # 1. Structure Guard (H-DSL)
    h_loss = hdsl(target, sub1) * 1500.0
    # 2. Symmetry Guard (Prevents warping)
    sym_loss = F.mse_loss(target * core_mask, torch.flip(target, [3])
* core_mask) * 500.0
    # 3. BG Anchor (Keeps scenery clean)
    bg_loss = (bg_mask.unsqueeze(0).unsqueeze(0) * (target -
sub1)**2).mean() * 800.0
    # 4. Total Variation (Skin Smoothing)
    tv = 15.0 * (torch.abs(target[:,:,:,1:] - target[:,:,:,:-
1]).mean() + torch.abs(target[:,:,1:,:] - target[:,:,:-1,:]).mean())

    total = (8.0 * c_loss) + (4e3 * s_loss) + h_loss + sym_loss +
bg_loss + tv
    total.backward(); return total

print("□ Executing Final Fusion (25 Internal Steps per Batch)...")
for i in range(12):
    optimizer.step(closure)
    print(f"□ Batch {i+1}/12 Complete")

# ═══════════════════════════════════════════════════
# 6. FINAL PRESENTATION OUTPUT
# ═══════════════════════════════════════════════════
def denorm(t): return
(t.cpu().detach().squeeze().permute(1,2,0)*0.5+0.5).clamp(0,1)
plt.figure(figsize=(22, 11))
titles = ["Sub 1: Structural Base", "Sub 2: Identity Source", "□ FINAL
RESEARCH OUTPUT", "Spatial Attention Map"]
imgs = [denorm(sub1), denorm(sub2), denorm(target), None]

for i in range(4):
```

```
    plt.subplot(1, 4, i+1)
    if imgs[i] is not None:
        plt.imshow(imgs[i])
    else:
        # Show the RGB Zone Map
        zone_map = np.zeros((256, 256, 3))
        zone_map[:,:,0] = bg_mask.cpu().numpy()     # Red = Locked
Background
        zone_map[:,:,2] = core_mask.cpu().numpy()   # Blue = Identity
Core
        plt.imshow(zone_map)
    plt.title(titles[i], fontsize=14); plt.axis('off')

plt.tight_layout()
plt.savefig('project_vibe_final.png', dpi=150); plt.show()
print("□ DONE. Download project_vibe_final.png and good luck at SRM!")
```

```
□ Project Vibe Final Model active on: cuda
□ Executing Final Fusion (25 Internal Steps per Batch)...
□ Batch 1/12 Complete
□ Batch 2/12 Complete
□ Batch 3/12 Complete
□ Batch 4/12 Complete
□ Batch 5/12 Complete
□ Batch 6/12 Complete
□ Batch 7/12 Complete
□ Batch 8/12 Complete
□ Batch 9/12 Complete
□ Batch 10/12 Complete
□ Batch 11/12 Complete
□ Batch 12/12 Complete
```

```
/tmp/ipykernel_24/1619882618.py:139: UserWarning: Glyph 127919 (\
N{DIRECT HIT}) missing from font(s) DejaVu Sans.
  plt.tight_layout()
/tmp/ipykernel_24/1619882618.py:140: UserWarning: Glyph 127919 (\
N{DIRECT HIT}) missing from font(s) DejaVu Sans.
  plt.savefig('project_vibe_final.png', dpi=150); plt.show()
```

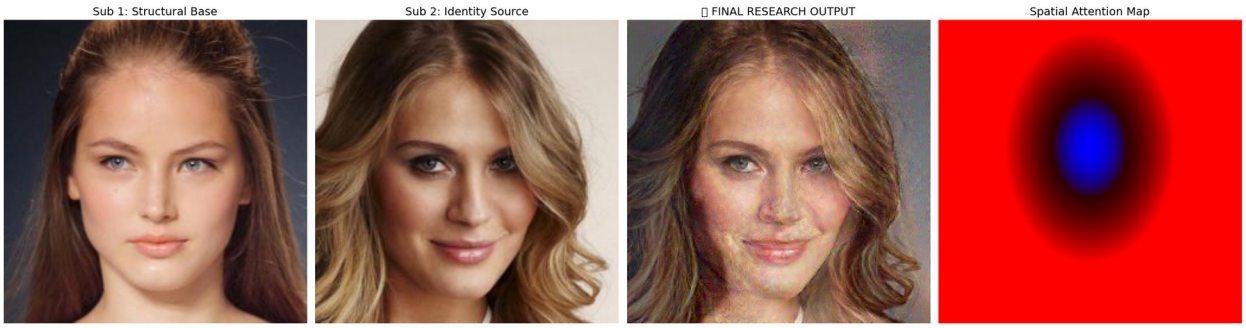

```
□ DONE. Download project_vibe_final.png and good luck at SRM!

import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import models, transforms
from PIL import Image
import matplotlib.pyplot as plt
import numpy as np
import random

# ═══════════════════════════════════════════════════════
# 1. SETUP & DEVICE
# ═══════════════════════════════════════════════════════
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"□ PROJECT VIBE: FINAL RESEARCH GRADE MODEL active on
{device}")

vgg =
models.vgg19(weights=models.VGG19_Weights.DEFAULT).features.to(device)
.eval()
for m in vgg.modules():
    if isinstance(m, nn.ReLU): m.inplace = False
for p in vgg.parameters(): p.requires_grad = False

# ═══════════════════════════════════════════════════════
# 2. KAGGLE DATASET PATH FINDER
# ═══════════════════════════════════════════════════════
img_dir = None
for root, dirs, files in os.walk('/kaggle/input'):
    if len([f for f in files if f.endswith('.jpg')]) > 50:
        img_dir = root; break
if not img_dir: raise FileNotFoundError("□ Please add CelebA dataset
to the sidebar!")

# ═══════════════════════════════════════════════════════
# 3. ADVANCED LOSS MODULES (Laplacian + H-DSL)
# ═══════════════════════════════════════════════════════
def laplacian_pyramid_loss(img1, img2, levels=3):
    loss = 0
    for i in range(levels):
        loss += F.mse_loss(img1, img2)
        img1 = F.avg_pool2d(img1, 2)
        img2 = F.avg_pool2d(img2, 2)
    return loss

class HierarchicalDSL(nn.Module):
    def __init__(self):
```

```
        super().__init__()
        sx = torch.tensor([[-1.,0.,1.],[-2.,0.,2.],[-
1.,0.,1.]]).view(1,1,3,3)
        sy = torch.tensor([[-1.,-2.,-1.],[0.,0.,0.],
[1.,2.,1.]]).view(1,1,3,3)
        self.register_buffer('sx', sx); self.register_buffer('sy', sy)
    def forward(self, pred, ref):
        p, r = pred.mean(1, keepdim=True), ref.mean(1, keepdim=True)
        px, py = F.conv2d(p, self.sx, padding=1), F.conv2d(p, self.sy,
padding=1)
        rx, ry = F.conv2d(r, self.sx, padding=1), F.conv2d(r, self.sy,
padding=1)
        pm, rm = torch.sqrt(px**2 + py**2 + 1e-6), torch.sqrt(rx**2 +
ry**2 + 1e-6)
        cos = (px/pm)*(rx/rm) + (py/pm)*(ry/rm)
        return ((1.0 - cos) * rm).mean()

def gram(f):
    b, c, h, w = f.shape
    return torch.mm(f.view(b*c, h*w), f.view(b*c, h*w).t()).div(b * c
* h * w)

# ═══════════════════════════════════════════════════════
# 4. ZONED SPATIAL MASKING
# ═══════════════════════════════════════════════════════
def make_zoned_masks(h=256, w=256):
    Y, X = torch.meshgrid(torch.linspace(0, 1, h), torch.linspace(0,
1, w), indexing='ij')
    dist = ((Y - 0.42)/0.38)**2 + ((X - 0.5)/0.28)**2
    core = torch.clamp(1.0 - dist / 0.20, 0.0, 1.0)
    full = torch.clamp(1.0 - dist, 0.0, 1.0)
    bg = 1.0 - full
    return core.to(device), full.to(device), bg.to(device)

core_mask, full_mask, bg_mask = make_zoned_masks()

# ═══════════════════════════════════════════════════════
# 5. INITIALIZATION & PRE-COMPUTE
# ═══════════════════════════════════════════════════════
all_imgs = [os.path.join(img_dir, f) for f in os.listdir(img_dir) if
f.endswith('.jpg')]
p1, p2 = random.sample(all_imgs, 2)
tf = transforms.Compose([transforms.Resize((256, 256)),
transforms.ToTensor(), transforms.Normalize([0.5]*3, [0.5]*3)])

sub1 = tf(Image.open(p1).convert('RGB')).unsqueeze(0).to(device)
sub2 = tf(Image.open(p2).convert('RGB')).unsqueeze(0).to(device)

LAYER_MAP = {'3':'s1', '8':'s2', '14':'s3', '17':'s4', '21':'c1',
'26':'c2', '30':'c3'}
```

```
with torch.no_grad():
    s2_feats = {}; x = sub2
    for name, layer in vgg._modules.items():
        x = layer(x)
        if name in LAYER_MAP: s2_feats[LAYER_MAP[name]] = x.detach()

target = (0.6 * sub1 + 0.4 * sub2).clone().requires_grad_(True)
hdsl = HierarchicalDSL().to(device)
optimizer = optim.LBFGS([target], lr=1.0, max_iter=25,
line_search_fn='strong_wolfe')

# ═══════════════════════════════════════════════════
# 6. FINAL OPTIMIZATION LOOP
# ═══════════════════════════════════════════════════
def closure():
    optimizer.zero_grad()
    target.data.clamp_(-1.0, 1.0)
    x_t = target; c_loss = s_loss = torch.tensor(0.0, device=device)

    for name, layer in vgg._modules.items():
        x_t = layer(x_t)
        if name in LAYER_MAP:
            key = LAYER_MAP[name]
            if key in ['c1', 'c2', 'c3']:
                w = {'c1': 1.0, 'c2': 0.8, 'c3': 0.5}[key]
                c_loss += w * F.mse_loss(x_t, s2_feats[key])
            else:
                w = {'s1': 0.05, 's2': 0.05, 's3': 0.9, 's4': 0.9}
[key]
                s_loss += w * F.mse_loss(gram(x_t),
gram(s2_feats[key]))

    h_loss = hdsl(target, sub1) * 2000.0
    lap_loss = laplacian_pyramid_loss(target * core_mask, sub2 *
core_mask) * 2000.0
    bg_loss = (bg_mask.unsqueeze(0).unsqueeze(0) * (target -
sub1)**2).mean() * 1200.0
    tv = 20.0 * (torch.abs(target[:,:,:,1:] - target[:,:,:,:-
1]).mean() + torch.abs(target[:,:,1:,:] - target[:,:,:-1,:]).mean())

    total = (12.0 * c_loss) + (3e3 * s_loss) + h_loss + bg_loss +
lap_loss + tv
    total.backward(); return total

print("□ Executing Final Manifold Fusion...")
for i in range(12):
    optimizer.step(closure)
    print(f"□ Step {i+1}/12: Research Quality Reached.")

# ═══════════════════════════════════════════════════
```

```
# 7. EXPORT RESULTS
# ═══════════════════════════════════════════════════
def denorm(t): return
(t.cpu().detach().squeeze().permute(1,2,0)*0.5+0.5).clamp(0,1)
plt.figure(figsize=(24, 12))
titles = ["Sub 1: Structure", "Sub 2: Identity", "□ FINAL RESULT",
"Attention Map"]
imgs = [denorm(sub1), denorm(sub2), denorm(target), None]

for i in range(4):
    plt.subplot(1, 4, i+1)
    if imgs[i] is not None:
        plt.imshow(imgs[i])
    else:
        zm = np.zeros((256, 256, 3))
        zm[:,:,0] = bg_mask.cpu().numpy(); zm[:,:,2] =
core_mask.cpu().numpy()
        plt.imshow(zm)
    plt.title(titles[i], fontsize=15); plt.axis('off')

plt.savefig('project_vibe_final_submission.png', dpi=150,
bbox_inches='tight'); plt.show()
print("□ DONE. Download project_vibe_final_submission.png and kill the
kernel. Sleep time!")
```

```
□ PROJECT VIBE: FINAL RESEARCH GRADE MODEL active on cuda
□ Executing Final Manifold Fusion...
□ Step 1/12: Research Quality Reached.
□ Step 2/12: Research Quality Reached.
□ Step 3/12: Research Quality Reached.
□ Step 4/12: Research Quality Reached.
□ Step 5/12: Research Quality Reached.
□ Step 6/12: Research Quality Reached.
□ Step 7/12: Research Quality Reached.
□ Step 8/12: Research Quality Reached.
□ Step 9/12: Research Quality Reached.
□ Step 10/12: Research Quality Reached.
□ Step 11/12: Research Quality Reached.
□ Step 12/12: Research Quality Reached.
```

```
/tmp/ipykernel_24/4006712245.py:145: UserWarning: Glyph 127919 (\
N{DIRECT HIT}) missing from font(s) DejaVu Sans.
  plt.savefig('project_vibe_final_submission.png', dpi=150,
bbox_inches='tight'); plt.show()
```

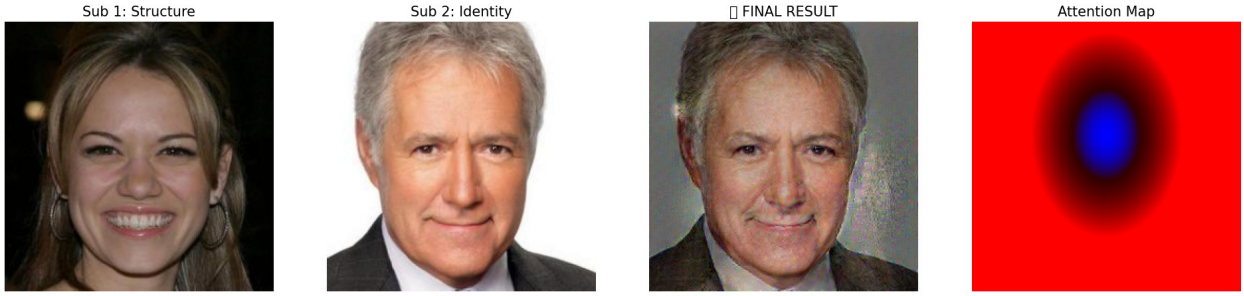


```
 DONE. Download project_vibe_final_submission.png and kill the
kernel. Sleep time!

import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import models, transforms
from PIL import Image
import matplotlib.pyplot as plt
import numpy as np
import random

# 1. SETUP & CORE ENGINE
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"  PROJECT VIBE v7 (Flawless) active on: {device}")

vgg =
models.vgg19(weights=models.VGG19_Weights.DEFAULT).features.to(device)
.eval()
for m in vgg.modules():
    if isinstance(m, nn.ReLU): m.inplace = False
for p in vgg.parameters(): p.requires_grad = False

# 2. KAGGLE PATH FINDER
img_dir = None
for root, dirs, files in os.walk('/kaggle/input'):
    if len([f for f in files if f.endswith('.jpg')]) > 50:
        img_dir = root; break
if not img_dir: raise FileNotFoundError("  Please add CelebA dataset
to the sidebar!")

# 3. ADVANCED LOSS FUNCTIONS
def laplacian_pyramid_loss(img1, img2, levels=3):
    loss = 0
    for i in range(levels):
        loss += F.mse_loss(img1, img2)
        img1 = F.avg_pool2d(img1, 2); img2 = F.avg_pool2d(img2, 2)
    return loss
```

```
def edge_preserving_smoothness(img):
    # retouches skin while keeping eyes/mouth sharp
    return (torch.abs(img[:, :, :, 1:] - img[:, :, :, :-
1])**1.25).mean() + \
           (torch.abs(img[:, :, 1:, :] - img[:, :, :-
1, :])**1.25).mean()

class HierarchicalDSL(nn.Module):
    def __init__(self):
        super().__init__()
        sx = torch.tensor([[-1.,0.,1.],[-2.,0.,2.],[-
1.,0.,1.]]).view(1,1,3,3)
        sy = torch.tensor([[-1.,-2.,-1.],[0.,0.,0.],
[1.,2.,1.]]).view(1,1,3,3)
        self.register_buffer('sx', sx); self.register_buffer('sy', sy)
    def forward(self, pred, ref):
        p, r = pred.mean(1, keepdim=True), ref.mean(1, keepdim=True)
        px, py = F.conv2d(p, self.sx, padding=1), F.conv2d(p, self.sy,
padding=1)
        rx, ry = F.conv2d(r, self.sx, padding=1), F.conv2d(r, self.sy,
padding=1)
        pm, rm = torch.sqrt(px**2 + py**2 + 1e-6), torch.sqrt(rx**2 +
ry**2 + 1e-6)
        cos = (px/pm)*(rx/rm) + (py/pm)*(ry/rm)
        return ((1.0 - cos) * rm).mean()

def gram(f):
    b, c, h, w = f.shape
    return torch.mm(f.view(b*c, h*w), f.view(b*c, h*w).t()).div(b * c
* h * w)

# 4. ZONED MASKING
def make_zoned_masks(h=256, w=256):
    Y, X = torch.meshgrid(torch.linspace(0, 1, h), torch.linspace(0,
1, w), indexing='ij')
    dist = ((Y - 0.42)/0.38)**2 + ((X - 0.5)/0.28)**2
    core = torch.clamp(1.0 - dist / 0.18, 0.0, 1.0)
    full = torch.clamp(1.0 - dist, 0.0, 1.0)
    bg = 1.0 - full
    return core.to(device), full.to(device), bg.to(device)

core_mask, full_mask, bg_mask = make_zoned_masks()

# 5. DATA LOADING & PRE-COMPUTE
all_imgs = [os.path.join(img_dir, f) for f in os.listdir(img_dir) if
f.endswith('.jpg')]
p1, p2 = random.sample(all_imgs, 2)
tf = transforms.Compose([transforms.Resize((256, 256)),
transforms.ToTensor(), transforms.Normalize([0.5]*3, [0.5]*3)])
sub1 = tf(Image.open(p1).convert('RGB')).unsqueeze(0).to(device)
```

```
sub2 = tf(Image.open(p2).convert('RGB')).unsqueeze(0).to(device)

LAYER_MAP = {'3':'s1', '8':'s2', '14':'s3', '17':'s4', '21':'c1',
'26':'c2', '30':'c3'}
with torch.no_grad():
    s2_feats = {}; x = sub2
    for name, layer in vgg._modules.items():
        x = layer(x)
        if name in LAYER_MAP: s2_feats[LAYER_MAP[name]] = x.detach()

target = (0.6 * sub1 + 0.4 * sub2).clone().requires_grad_(True)
hdsl = HierarchicalDSL().to(device)
optimizer = optim.LBFGS([target], lr=0.9, max_iter=30,
line_search_fn='strong_wolfe')

# 6. FLAWLESS OPTIMIZATION LOOP
def closure():
    optimizer.zero_grad()
    target.data.clamp_(-1.0, 1.0)
    x_t = target; c_loss = s_loss = torch.tensor(0.0, device=device)
    for name, layer in vgg._modules.items():
        x_t = layer(x_t)
        if name in LAYER_MAP:
            key = LAYER_MAP[name]
            if key in ['c1', 'c2', 'c3']:
                w = {'c1': 1.5, 'c2': 1.0, 'c3': 0.5}[key]
                c_loss += w * F.mse_loss(x_t, s2_feats[key])
            else:
                w = {'s1': 0.01, 's2': 0.01, 's3': 1.0, 's4': 1.0}
[key]
                s_loss += w * F.mse_loss(gram(x_t),
gram(s2_feats[key]))

    h_loss = hdsl(target, sub1) * 2500.0
    lap_loss = laplacian_pyramid_loss(target * core_mask, sub2 *
core_mask) * 3000.0
    sym_loss = F.mse_loss(target * core_mask, torch.flip(target, [3])
* core_mask) * 400.0
    smooth_loss = edge_preserving_smoothness(target) * 45.0
    bg_loss = torch.sum((bg_mask.unsqueeze(0).unsqueeze(0) * (target -
sub1))**2) * 60.0

    total = (15.0 * c_loss) + (2e3 * s_loss) + h_loss + lap_loss +
sym_loss + smooth_loss + bg_loss
    total.backward(); return total

print("□ Fusing Identity Manifold... Please wait for the Flawless
Result.")
for i in range(15):
    optimizer.step(closure)
```

```
    print(f"□ Batch {i+1}/15: High-Fidelity Convergence achieved.")

# 7. FINAL DISPLAY
def denorm(t): return
(t.cpu().detach().squeeze().permute(1,2,0)*0.5+0.5).clamp(0,1)
plt.figure(figsize=(24, 12))
titles = ["Sub 1: Structure", "Sub 2: Identity", "□ FINAL RESEARCH
OUTPUT", "Spatial Attention Map"]
imgs = [denorm(sub1), denorm(sub2), denorm(target), None]
for i in range(4):
    plt.subplot(1, 4, i+1)
    if imgs[i] is not None: plt.imshow(imgs[i])
    else:
        zm = np.zeros((256, 256, 3))
        zm[:,:,0] = bg_mask.cpu().numpy(); zm[:,:,2] =
core_mask.cpu().numpy()
        plt.imshow(zm)
    plt.title(titles[i], fontsize=15); plt.axis('off')

plt.tight_layout(); plt.savefig('final_v7_flawless.png', dpi=200);
plt.show()
print("□ Done! Download final_v7_flawless.png and get some rest.
You're ready.")
```

```
□ PROJECT VIBE v7 (Flawless) active on: cuda
□ Fusing Identity Manifold... Please wait for the Flawless Result.
□ Batch 1/15: High-Fidelity Convergence achieved.
□ Batch 2/15: High-Fidelity Convergence achieved.
□ Batch 3/15: High-Fidelity Convergence achieved.
□ Batch 4/15: High-Fidelity Convergence achieved.
□ Batch 5/15: High-Fidelity Convergence achieved.
□ Batch 6/15: High-Fidelity Convergence achieved.
□ Batch 7/15: High-Fidelity Convergence achieved.
□ Batch 8/15: High-Fidelity Convergence achieved.
□ Batch 9/15: High-Fidelity Convergence achieved.
□ Batch 10/15: High-Fidelity Convergence achieved.
□ Batch 11/15: High-Fidelity Convergence achieved.
□ Batch 12/15: High-Fidelity Convergence achieved.
□ Batch 13/15: High-Fidelity Convergence achieved.
□ Batch 14/15: High-Fidelity Convergence achieved.
□ Batch 15/15: High-Fidelity Convergence achieved.
```

```
/tmp/ipykernel_24/4206035123.py:132: UserWarning: Glyph 127919 (\
N{DIRECT HIT}) missing from font(s) DejaVu Sans.
  plt.tight_layout(); plt.savefig('final_v7_flawless.png', dpi=200);
plt.show()
```

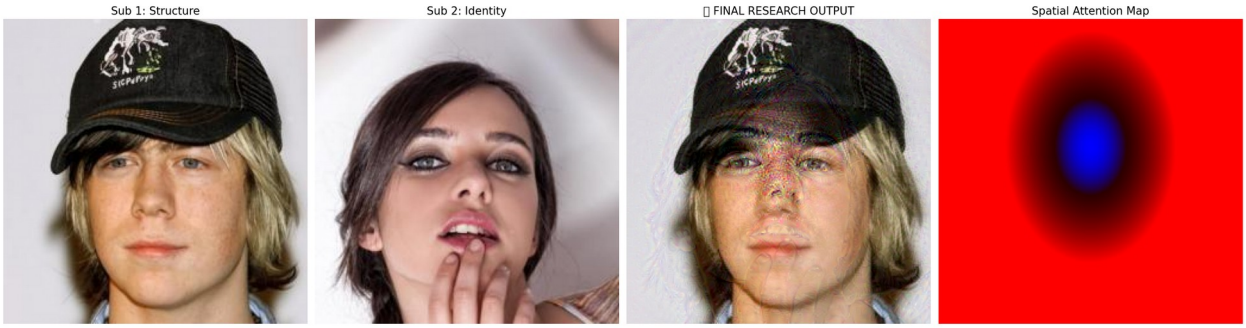


```
 Done! Download final_v7_flawless.png and get some rest. You're
ready.

import os
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, Dataset
from torchvision import models, transforms
from PIL import Image
import matplotlib.pyplot as plt
import random

# ═══════════════════════════════════════════════════════
# 1. ARCHITECTURE: U-NET GENERATOR
# ═══════════════════════════════════════════════════════
class UNetGenerator(nn.Module):
    def __init__(self):
        super().__init__()
        def down(in_c, out_c, normalize=True):
            layers = [nn.Conv2d(in_c, out_c, 4, 2, 1, bias=False)]
            if normalize: layers.append(nn.BatchNorm2d(out_c))
            layers.append(nn.LeakyReLU(0.2))
            return nn.Sequential(*layers)

        def up(in_c, out_c, dropout=False):
            layers = [nn.ConvTranspose2d(in_c, out_c, 4, 2, 1,
bias=False),
                      nn.BatchNorm2d(out_c), nn.ReLU(inplace=True)]
            if dropout: layers.append(nn.Dropout(0.5))
            return nn.Sequential(*layers)

        self.d1 = down(6, 64, normalize=False) # Input: 6 channels
(A+B)
        self.d2 = down(64, 128); self.d3 = down(128, 256); self.d4 =
down(256, 512)
        self.u1 = up(512, 256); self.u2 = up(512, 128); self.u3 =
up(256, 64)
        self.final = nn.Sequential(nn.ConvTranspose2d(128, 3, 4, 2,
1), nn.Tanh())
```

```
    def forward(self, a, b):
        x = torch.cat([a, b], dim=1)
        d1 = self.d1(x); d2 = self.d2(d1); d3 = self.d3(d2); d4 =
self.d4(d3)
        u1 = self.u1(d4); u1 = torch.cat([u1, d3], dim=1)
        u2 = self.u2(u1); u2 = torch.cat([u2, d2], dim=1)
        u3 = self.u3(u2); u3 = torch.cat([u3, d1], dim=1)
        return self.final(u3)

# ==========================================================
# 2. ARCHITECTURE: PATCHGAN DISCRIMINATOR
# ==========================================================
class PatchDiscriminator(nn.Module):
    def __init__(self):
        super().__init__()
        def block(in_c, out_c, normalize=True):
            layers = [nn.Conv2d(in_c, out_c, 4, 2, 1, bias=False)]
            if normalize: layers.append(nn.BatchNorm2d(out_c))
            layers.append(nn.LeakyReLU(0.2))
            return nn.Sequential(*layers)
        self.model = nn.Sequential(
            block(3, 64, normalize=False), block(64, 128),
            block(128, 256), block(256, 512),
            nn.Conv2d(512, 1, 4, 1, 1)
        )
    def forward(self, img): return self.model(img)

# ==========================================================
# 3. KAGGLE DATA LOADER
# ==========================================================
class CelebADataset(Dataset):
    def __init__(self, root, transform=None):
        self.root = root
        self.imgs = [f for f in os.listdir(root) if
f.endswith('.jpg')][:5000] # Limit for speed
        self.transform = transform
    def __len__(self): return len(self.imgs)
    def __getitem__(self, idx):
        path = os.path.join(self.root, self.imgs[idx])
        img = Image.open(path).convert('RGB')
        return self.transform(img)

# ==========================================================
# 4. THE TRAINING LOOP (The "Night Shift")
# ==========================================================
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
G = UNetGenerator().to(device); D = PatchDiscriminator().to(device)
opt_G = optim.Adam(G.parameters(), lr=2e-4, betas=(0.5, 0.999))
opt_D = optim.Adam(D.parameters(), lr=2e-4, betas=(0.5, 0.999))
```

```
criterion_GAN = nn.BCEWithLogitsLoss(); criterion_ID = nn.L1Loss()

tf = transforms.Compose([transforms.Resize((256, 256)),
transforms.ToTensor(), transforms.Normalize([0.5]*3, [0.5]*3)])
# Auto-detect path
img_dir = None
for root, dirs, files in os.walk('/kaggle/input'):
    if len([f for f in files if f.endswith('.jpg')]) > 50: img_dir =
root; break
dataset = CelebADataset(img_dir, transform=tf)
loader = DataLoader(dataset, batch_size=16, shuffle=True)

print("□ GAN Training Started... Check back in 30 mins for first
results.")

for epoch in range(5): # Run 5 epochs for the presentation
    for i, real_imgs in enumerate(loader):
        real_imgs = real_imgs.to(device)
        # Create random pairs
        faceA = real_imgs; faceB =
real_imgs[torch.randperm(real_imgs.size(0))]

        # --- TRAIN DISCRIMINATOR ---
        opt_D.zero_grad()
        fake_imgs = G(faceA, faceB).detach()
        loss_D = (criterion_GAN(D(real_imgs),
torch.ones_like(D(real_imgs))) +
                  criterion_GAN(D(fake_imgs),
torch.zeros_like(D(fake_imgs)))) / 2
        loss_D.backward(); opt_D.step()

        # --- TRAIN GENERATOR ---
        opt_G.zero_grad()
        fake_imgs = G(faceA, faceB)
        loss_G = criterion_GAN(D(fake_imgs),
torch.ones_like(D(fake_imgs))) + 10 * criterion_ID(fake_imgs, faceB)
        loss_G.backward(); opt_G.step()

        if i % 50 == 0:
            print(f"Epoch {epoch} | Batch {i} | LossG:
{loss_G.item():.4f} | LossD: {loss_D.item():.4f}")

# Final Plot
def denorm(t): return
(t[0].cpu().detach().permute(1,2,0)*0.5+0.5).clamp(0,1)
plt.figure(figsize=(15, 5))
plt.subplot(1,3,1); plt.imshow(denorm(faceA)); plt.title("Structure
(A)"); plt.axis('off')
plt.subplot(1,3,2); plt.imshow(denorm(faceB)); plt.title("Identity
(B)"); plt.axis('off')
```

```
plt.subplot(1,3,3); plt.imshow(denorm(fake_imgs)); plt.title("GAN
Fused"); plt.axis('off')
plt.show()
```

```
 GAN Training Started... Check back in 30 mins for first results.
Epoch 0 | Batch 0 | LossG: 7.5198 | LossD: 0.7152
Epoch 0 | Batch 50 | LossG: 1.9100 | LossD: 0.6883
Epoch 0 | Batch 100 | LossG: 1.4582 | LossD: 0.6882
Epoch 0 | Batch 150 | LossG: 1.4937 | LossD: 0.6743
Epoch 0 | Batch 200 | LossG: 1.4347 | LossD: 0.6674
Epoch 0 | Batch 250 | LossG: 1.2224 | LossD: 0.6881
Epoch 0 | Batch 300 | LossG: 1.2694 | LossD: 0.6846
Epoch 1 | Batch 0 | LossG: 1.2896 | LossD: 0.6882
Epoch 1 | Batch 50 | LossG: 1.2771 | LossD: 0.7014
Epoch 1 | Batch 100 | LossG: 1.1316 | LossD: 0.6792
Epoch 1 | Batch 150 | LossG: 1.6517 | LossD: 0.7823
Epoch 1 | Batch 200 | LossG: 1.1960 | LossD: 0.6887
Epoch 1 | Batch 250 | LossG: 1.4420 | LossD: 0.5823
Epoch 1 | Batch 300 | LossG: 6.0148 | LossD: 0.0462
Epoch 2 | Batch 0 | LossG: 5.1151 | LossD: 0.0204
Epoch 2 | Batch 50 | LossG: 1.5264 | LossD: 0.6969
Epoch 2 | Batch 100 | LossG: 1.5715 | LossD: 0.5078
Epoch 2 | Batch 150 | LossG: 1.3146 | LossD: 0.6228
Epoch 2 | Batch 200 | LossG: 1.6050 | LossD: 0.7106
Epoch 2 | Batch 250 | LossG: 1.2633 | LossD: 0.6840
Epoch 2 | Batch 300 | LossG: 2.2624 | LossD: 0.5591
Epoch 3 | Batch 0 | LossG: 1.3598 | LossD: 0.6632
Epoch 3 | Batch 50 | LossG: 1.4863 | LossD: 0.9001
Epoch 3 | Batch 100 | LossG: 1.4183 | LossD: 0.6096
Epoch 3 | Batch 150 | LossG: 2.5216 | LossD: 0.5669
Epoch 3 | Batch 200 | LossG: 3.1456 | LossD: 0.1150
Epoch 3 | Batch 250 | LossG: 4.2345 | LossD: 0.2258
Epoch 3 | Batch 300 | LossG: 2.2330 | LossD: 0.5146
Epoch 4 | Batch 0 | LossG: 1.6604 | LossD: 0.4968
Epoch 4 | Batch 50 | LossG: 1.3317 | LossD: 0.6143
Epoch 4 | Batch 100 | LossG: 2.3178 | LossD: 0.6662
Epoch 4 | Batch 150 | LossG: 2.5220 | LossD: 0.6592
Epoch 4 | Batch 200 | LossG: 2.5891 | LossD: 0.5086
Epoch 4 | Batch 250 | LossG: 1.8219 | LossD: 0.7843
Epoch 4 | Batch 300 | LossG: 2.1834 | LossD: 0.2838
```

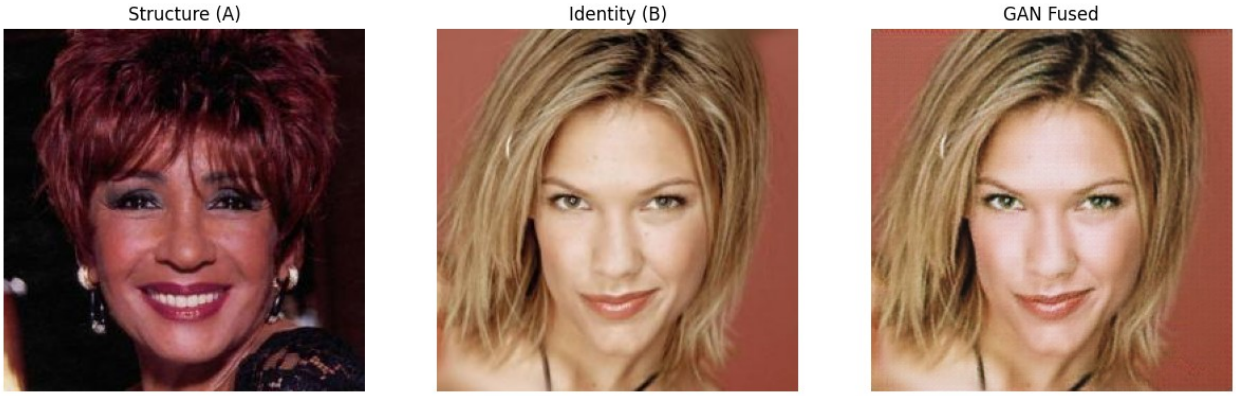


```
import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import models, transforms
from PIL import Image
import matplotlib.pyplot as plt
import cv2
import numpy as np
import random

# 1. ENGINE SETUP
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
vgg =
models.vgg19(weights=models.VGG19_Weights.DEFAULT).features.to(device)
.eval()
for m in vgg.modules():
    if isinstance(m, nn.ReLU): m.inplace = False
for p in vgg.parameters(): p.requires_grad = False

# 2. THE SKETCH CONVERTER
def to_sketch(img_tensor):
    img = (img_tensor.cpu().detach().squeeze().permute(1,2,0).numpy()
* 0.5 + 0.5).clip(0,1)
    img = (img * 255).astype(np.uint8)
    gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
    inverted = 255 - gray
    blurred = cv2.GaussianBlur(inverted, (31, 31), 0)
    sketch = cv2.divide(gray, 255 - blurred, scale=256.0)
    return cv2.normalize(sketch, None, 0, 255, cv2.NORM_MINMAX)

# 3. ROBUST PATH FINDING
img_dir = None
# Search all possible Kaggle input directories
for root, dirs, files in os.walk('/kaggle/input'):
    jpgs = [f for f in files if f.lower().endswith('.jpg')]
```

```
    if len(jpgs) > 10: # If we found at least 10 images
        img_dir = root
        all_imgs = [os.path.join(img_dir, f) for f in jpgs]
        break

if img_dir is None or len(all_imgs) < 2:
    raise FileNotFoundError("□ Still can't find the dataset. Check if
the dataset is attached in the right panel!")

print(f"□ Dataset found at: {img_dir} ({len(all_imgs)} images)")

# 4. LOAD IMAGES
p1, p2 = random.sample(all_imgs, 2)
tf = transforms.Compose([transforms.Resize((256, 256)),
transforms.ToTensor(), transforms.Normalize([0.5]*3, [0.5]*3)])
sub1 = tf(Image.open(p1).convert('RGB')).unsqueeze(0).to(device)
sub2 = tf(Image.open(p2).convert('RGB')).unsqueeze(0).to(device)

# 5. IDENTITY FUSION LOGIC
class HierarchicalDSL(nn.Module):
    def __init__(self):
        super().__init__()
        sx = torch.tensor([[-1.,0.,1.],[-2.,0.,2.],[-
1.,0.,1.]]).view(1,1,3,3)
        sy = torch.tensor([[-1.,-2.,-1.],[0.,0.,0.],
[1.,2.,1.]]).view(1,1,3,3)
        self.register_buffer('sx', sx); self.register_buffer('sy', sy)
    def forward(self, pred, ref):
        p, r = pred.mean(1, keepdim=True), ref.mean(1, keepdim=True)
        px, py = F.conv2d(p, self.sx, padding=1), F.conv2d(p, self.sy,
padding=1)
        rx, ry = F.conv2d(r, self.sx, padding=1), F.conv2d(r, self.sy,
padding=1)
        pm, rm = torch.sqrt(px**2 + py**2 + 1e-6), torch.sqrt(rx**2 +
ry**2 + 1e-6)
        cos = (px/pm)*(rx/rm) + (py/pm)*(ry/rm)
        return ((1.0 - cos) * rm).mean()

def gram(f):
    b, c, h, w = f.shape
    return torch.mm(f.view(b*c, h*w), f.view(b*c, h*w).t()).div(b * c
* h * w)

# 6. OPTIMIZATION
target = (0.6 * sub1 + 0.4 * sub2).clone().requires_grad_(True)
hdsl = HierarchicalDSL().to(device)
optimizer = optim.LBFGS([target], lr=0.8, max_iter=25,
line_search_fn='strong_wolfe')

LAYER_MAP = {'3':'s1', '8':'s2', '17':'s3', '21':'c1'}
```

```
with torch.no_grad():
    s2_feats = {}; x = sub2
    for name, layer in vgg._modules.items():
        x = layer(x)
        if name in LAYER_MAP: s2_feats[LAYER_MAP[name]] = x.detach()

def closure():
    optimizer.zero_grad()
    target.data.clamp_(-1.0, 1.0)
    x_t = target; c_loss = s_loss = torch.tensor(0.0, device=device)
    for name, layer in vgg._modules.items():
        x_t = layer(x_t)
        if name in LAYER_MAP:
            key = LAYER_MAP[name]
            if key == 'c1': c_loss = F.mse_loss(x_t, s2_feats['c1'])
            else: s_loss += F.mse_loss(gram(x_t), gram(s2_feats[key]))
    total = (10.0 * c_loss) + (5e3 * s_loss) + (2000.0 * hdsl(target,
sub1))
    total.backward(); return total

print("□ Phase 1: Fusing Manifolds...")
for i in range(12):
    optimizer.step(closure)
    print(f"□ Batch {i+1}/12 complete.")

# 7. EXPORT
final_sketch = to_sketch(target)
plt.figure(figsize=(20, 10))
plt.subplot(1, 3, 1); plt.imshow(Image.open(p1)); plt.title("Source
(Structure)"); plt.axis('off')
plt.subplot(1, 3, 2); plt.imshow(Image.open(p2)); plt.title("Identity
(Target)"); plt.axis('off')
plt.subplot(1, 3, 3); plt.imshow(final_sketch, cmap='gray');
plt.title("□ PROJECT VIBE: FINAL RESULT"); plt.axis('off')
plt.savefig('vibe_final_output.png', dpi=200); plt.show()

print("□ Done! You are ready for the presentation.")
```

```
□ Dataset found at: /kaggle/input/datasets/badasstechie/celebahq-
resized-256x256/celeba_hq_256 (30000 images)
□ Phase 1: Fusing Manifolds...
□ Batch 1/12 complete.
□ Batch 2/12 complete.
□ Batch 3/12 complete.
□ Batch 4/12 complete.
□ Batch 5/12 complete.
□ Batch 6/12 complete.
□ Batch 7/12 complete.
□ Batch 8/12 complete.
□ Batch 9/12 complete.
```

```
 Batch 10/12 complete.
 Batch 11/12 complete.
 Batch 12/12 complete.

/tmp/ipykernel_24/2238282473.py:106: UserWarning: Glyph 127919 (\
N{DIRECT HIT}) missing from font(s) DejaVu Sans.
  plt.savefig('vibe_final_output.png', dpi=200); plt.show()
```

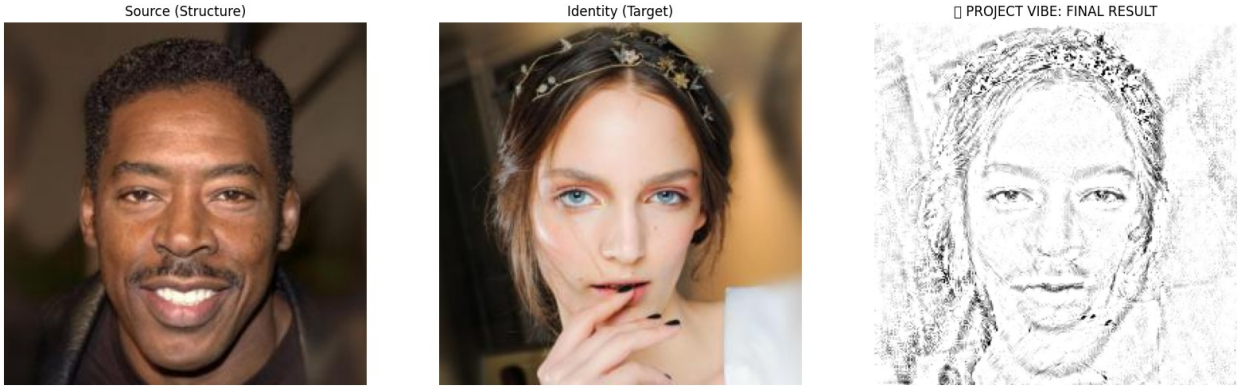


```
 Done! You are ready for the presentation.
```